Memory Is Not Enough: Separating Drift from Background Conflict in Continual Document Understanding

Supplementary Material

A. Setup details

DIL unified schema. A single four-entity label space {HEADER, KEY, VALUE, OTHER} (9 BIO tags) is held fixed over FUNSD→SROIE→CORD. FUNSD: HEADER→HEADER, QUESTION→KEY, ANSWER→VALUE. SROIE: COMPANY→KEY; ADDRESS/DATE/TOTAL→VALUE. CORD: menu/sub_total/total→VALUE; void_menu/sub→OTHER. KEY originates only from FUNSD and SROIE, so the schema is VALUE-dominant; we report per-class F1 for every DIL result and treat a high aggregate as retention only if the sparse classes are retained. The reverse order is used as a robustness check.

CIL-CORD sessions. CORD's 30 fine-grained classes in five sessions of six, ordered easy-to-hard by superclass: (0) core menu fields; (1) extended menu + sub; (2) sub_total fields; (3) total / payment methods; (4) residual total + void/auxiliary. Every receipt containing at least one in-session entity is kept, with out-of-session tokens relabelled to "O", giving per-session training pools of 800/382/549/723/780 receipts. Class imbalance across sessions is deliberately left unmodified. Joint uses the full-label pool (each receipt once with all its labels) rather than the concatenation of the masked per-session copies, which would feed conflicting labels for the same document.

Mixed. Six sessions: CORD-menu → FUNSD → CORD-menu-extended → SROIE → CORD-sub_total → FUNSD-revisit; boundaries alternate new-class and new-domain transitions and the final session is an explicit retention check.

Datasets. FUNSD (`nielsr/funsd-layoutlmv3`; 149/50), CORD (`naver-clova-ix/cord-v2`; 800/100/100), SROIE (`mp-02/sroie`; 626/347). Dataset-provided OCR and boxes throughout; images 224×224; sequence length 512.

Hyperparameters. Every run uses AdamW, lr 5×10^{-5} constant, weight decay 0.01 and clipping 1.0. EWC λ=10^3, Fisher over 200 documents; LwF T=2, α=1; ER/DER++ reservoir of 200 samples, DER++ α=β=0.5. Our O-LoRA and CL-LoRA ports train a single rank-8 LoRA adapter on the query, key and value projections of the frozen pretrained encoder, shared by all sessions and retrained in each with an orthogonality penalty (λ=0.5) against its own earlier states; CL-LoRA adds distillation towards the previous model (T=2). The published methods keep one adapter per task, so these rows report our ports, not the published designs. Early stopping uses patience 2 and δ=0.1 F1 points. Seeds 42/7/123. The epoch cap, numeric precision, early-stopping split and memory size differ between groups of legacy-protocol (†) cells; Tab. 7 lists them, and the clean protocol is specified in Sec. 3 of the main paper. Image CIL (CIFAR-100, ImageNet-R; ViT-B/16, 10 sessions): every reported cell, including the frozen-backbone and replay baselines, was trained on a 6 GB GPU at batch size 16 in fp32 with gradient checkpointing, 100-epoch cap with early stopping (patience 2), seeds 42/7/123; the new CODA-Prompt, CL-LoRA and LCA image cells use this recipe unchanged.

Table 1. Prompt/LoRA/C-Flat++ baselines, DIL, AA and BWT, mean±std over seeds 42/7/123. Naive 41.3 (LayoutLMv3), ER 87.9.

	LayoutLMv3		BERT	
Method	AA	BWT	AA	BWT
L2P	30.4 ± 0.4	-11.5 ± 0.1	34.4 ± 0.7	-38.5 ± 1.0
DualPrompt	29.9 ± 0.6	-11.1 ± 0.2	34.0 ± 0.2	-40.2 ± 1.1
CODA-Prompt	32.5 ± 1.2	-10.5 ± 0.2	35.2 ± 0.7	-40.7 ± 1.3
O-LoRA	40.8 ± 0.3	-56.8 ± 1.0	37.4 ± 0.3	-53.4 ± 0.9
CL-LoRA	41.1 ± 0.3	-58.8 ± 2.1	38.6 ± 0.4	-31.6 ± 1.9
ER + C-Flat++	88.4 ± 0.1	-1.8 ± 0.6	78.3 ± 0.4	-2.2 ± 0.1

Compute. Every run fits a single GPU; the full generality grid is roughly 200 GPU-hours on a 24 GB card.

B. Additional results

Prompt-, LoRA- and flatness-based baselines. Tab. 1 reports the frozen-backbone prompt family (L2P, DualPrompt, CODA-Prompt), the orthogonal/continual LoRA family and ER with C-Flat++ sharpness-aware optimisation on DIL, three seeds, LayoutLMv3 and BERT. The prompt family forgets little but never acquires the task (AA ≈ 30–35, below the naive level of 41.3): a frozen document encoder with a drifting shared head cannot be routed back into place. Our single-adapter LoRA ports acquire but forget as much as naive training does. Only the replay-based variant reaches the level of replay, and its sharpness term adds nothing over plain ER (88.4 vs. 87.9 on LayoutLMv3), consistent with Sec. 4.1 of the main paper: the remedy ordering is set by whether the head is re-grounded, not by the optimiser.

Table 2. Image class-incremental learning, AA and BWT, mean±std over seeds 42/7/123 (ViT-B/16, ten sessions). Fixed comparison set first, our rows, joint training (reference), then the dominated rows kept as evidence. SimpleCIL is a single seed.

Method	CIFAR-100 AA	CIFAR-100 BWT	ImageNet-R AA	ImageNet-R BWT
Naive	11.7 ± 1.1	-95.6 ± 1.4	9.3 ± 0.3	-87.1 ± 0.8
LwF	15.7 ± 0.6	-91.0 ± 0.7	12.7 ± 1.4	-84.4 ± 1.4
ER (2,000 exemplars)	76.2 ± 1.2	-23.8 ± 1.3	54.8 ± 0.7	-37.9 ± 0.9
DER++ (2,000 exemplars)	76.5 ± 1.5	-23.6 ± 1.7	52.8 ± 1.1	-40.5 ± 1.5
CODA-Prompt	15.2 ± 1.8	-89.7 ± 1.9	8.4 ± 0.7	-67.4 ± 1.7
CL-LoRA (2025)	14.1 ± 1.2	-92.7 ± 1.4	11.4 ± 0.8	-82.2 ± 0.8
SLCA (2023)	87.5 ± 0.2	-7.8 ± 0.2	65.0 ± 0.3	-13.5 ± 0.2
RanPAC (2023)	85.3 ± 0.1	-5.0 ± 0.1	58.5 ± 0.3	-7.0 ± 0.3
LCA (2026)	77.2 ± 3.2	-15.8 ± 3.8	54.7 ± 3.6	-21.8 ± 4.6
Latent memory, 500/class (ours, uncompressed)	80.0 ± 0.7	-19.8 ± 0.7	55.1 ± 1.3	-37.4 ± 1.3
CoLaR r=128, 500/class (ours)	79.5 ± 1.0	-20.4 ± 1.1	55.3 ± 1.4	-36.8 ± 1.5
Joint training (reference)	89.5 ± 0.4	–	77.6 ± 0.5	–
EWC	13.9 ± 1.5	-93.1 ± 1.6	9.3 ± 0.2	-87.2 ± 0.1
ER (200 exemplars)	42.3 ± 2.7	-61.6 ± 3.1	21.6 ± 1.3	-75.2 ± 1.4
DER++ (200 exemplars)	43.6 ± 4.0	-60.6 ± 4.2	20.2 ± 2.5	-77.2 ± 2.5
SLCA without classifier alignment	38.0 ± 3.9	-66.7 ± 4.3	25.8 ± 0.7	-67.4 ± 0.9
SimpleCIL (frozen trunk, NCM)	81.7 (1 seed)	-5.5 (1 seed)	48.2 (1 seed)	-6.2 (1 seed)
Per-class Gaussian latent replay (falsified)	10.7 ± 0.8	-96.5 ± 1.0	–	–

Image CIL: forgetting and the dominated rows. Tab. 2 extends the image-side comparison referenced in Sec. 5.1 of the main paper with BWT and with every image run outside the frozen reference set that we keep as evidence, all under the same recipe (ViT-B/16, ten sessions, batch 16, early stopping, seeds 42/7/123; cells not yet on disk are marked pending, never imputed). Two readings matter. Replay at 200 exemplars is far below the 2,000-exemplar rows (42.3/21.6 vs. 76.2/54.8 for ER), so the memory size in Tab. 4 is the one at which replay is competitive, not a favourable choice for our rows. SLCA's whole margin comes from its classifier alignment: without it the same recipe falls to 38.0/25.8, below 200-exemplar replay on CIFAR-100, which is the image-side instance of the head-locus finding. Per-class Gaussian latent replay, the image port of the document-side Gaussian store, sits at the naive level (10.7).

Ridge decomposition of the plastic-head gap. Tab. 3 reports a ridge-readout probe on the trained trunk of our replay rows (seed 42). After the last session the trunk is frozen and a closed-form ridge readout (random projection to 10,000 ReLU features, λ=10, the RanPAC recipe) is fitted under three data conditions: on the stored memory only (deployable), on all training data of all sessions (a non-deployable ceiling on what the features support), and on all data through the frozen pre-trained trunk (image only). Three differences follow: the *head gap* (full ridge minus the trained head) is the total headroom of the readout; the *coverage gap* (full minus stored) is the part that needs features of inputs the memory does not hold; the *trunk gap* (adapted minus frozen) is what replay-driven adaptation bought. On images the head gap is 9.5/16.1 points and is mostly coverage (7.5/14.3), so a better analytic head on the stored latents does not close the gap to SLCA; the adapted trunk is nevertheless 2.9/15.0 points better than the frozen one under the same readout, so adaptation is not the problem. On documents the head gap of CoLaR is within one point on DIL and Mixed and 7.7 points on CIL-CORD with zero coverage gap, i.e. the compressed store already covers its documents and the residual is the plastic head itself. The probe reads every past training image once at the end and is therefore a diagnostic, not a continual method.

Table 3. Ridge decomposition, seed 42, AA (LayoutLMv3 on documents, ViT-B/16 on images). *Head*: the trained plastic head. *Ridge, stored/full/frozen*: closed-form readout on the stored latents, on all seen training data, and on all data through the frozen pre-trained trunk. Gaps as defined in the text; a dash means the condition does not apply (documents have no frozen-trunk reference).

Cell	Head	Ridge, stored	Ridge, full	Ridge, frozen	Head gap	Trunk gap	Coverage gap
CIFAR-100, latent memory	78.8	80.8	88.3	85.4	9.5	2.9	7.5
ImageNet-R, latent memory	56.2	58.0	72.3	57.3	16.1	15.0	14.3
DIL, CoLaR	90.9	91.1	91.3	–	0.3	–	0.2
CIL-CORD, CoLaR	67.5	75.2	75.2	–	7.7	–	0.0
Mixed, CoLaR	70.9	71.0	71.6	–	0.7	–	0.6

Per-backbone forgetting. The full BWT and AF tables for every backbone×scenario×method cell of the generality grid are generated from the run directories and will ship with the code; Tab. 8 carries the AA view.

Budget-adaptive allocation is a frontier, not a gain. Allocating a fixed compressed-replay budget unevenly across tasks (documents $[18, 5, 2]$, ranks $[64, 128, 64]$ vs. a uniform 10 documents at rank 64; 5 epochs, k=4) gave +4.0 AA on seed 42 while passing a per-domain guard fixed in advance. On seeds 7 and 123 the same comparison gave −8.0 and +5.7: $+0.6 \pm 7.5$ over three seeds, a null. The seed-42 result was a point on a Pareto frontier (old-domain retention traded against the final domain at ≈ 9:1), not a shift of the frontier, and we report it as such.

Null cells among the fixed baselines (CPFD, IS3, CODA-Prompt). Three of the fixed baselines do not converge on the document scenarios, and we report those cells as nulls rather than as scores. CPFD and IS3 were ported onto the single growing token-classification head used by every other method here (the reference implementations keep per-task classifiers). The CIL-CORD and Mixed values for CPFD and IS3 below come from the earlier grid (fp32, 10-epoch cap, test-split early stopping, legacy scoring; Tab. 9); the session-restricted values in the main paper's comparison table come from later re-runs with a 100-epoch cap, which are different models, not a rescoring of these. On CIL-CORD, CPFD learns the first session (90–95 F1) and then reads 0.00 on every session from the first class-adding boundary (AA 2.7 and 0.0 on seeds 42 and 123); on Mixed it reads 6.0, 0.0 and 0.1; and on DIL, where no class is ever added,

the first domain falls to 8.0/4.0/0.9 F1 at the first boundary and the final row is zero on every domain for seed 42 and 3.0/9.7/0.0 for seed 7 (AA 0.0/4.2/14.2, reported as 6.1 ± 7.3 since one seed does learn the last domain), so the collapse is not a property of head growth. Two single-knob ablations on seed 42 reproduce the collapse exactly: flooring the per-sample CE weight at 0.5 (session-1 F1 0.00) and switching the distillation term off entirely (session-1 F1 2.1), so neither the adaptive CE factor nor the attention/logit distillation is the cause; what remains is the pseudo-labelling and O-token masking path, which removes the training signal on the new session's entities. IS3 learns each new session (diagonal 93/29/32/66/10 on seed 123) but never predicts a previously learned class again (every off-diagonal entry is exactly 0.00), giving AA 4.4 ± 2.2 on CIL-CORD in the earlier grid (three seeds); on Mixed, whose last session adds no classes, it recovers there to 30.1 ± 1.2. The main comparison reports the session-restricted re-runs instead. On DIL the nine-tag head never grows, so IS3's old/new split is degenerate (every entity row is "old" at every boundary and the O-token distillation spans the whole output) and the method reduces to sequential fine-tuning under a full-output distillation term; seed 42 learns the first domain (88.2) and drops both earlier domains below 8 F1 at the next boundary (AA 38.3 vs. naive 41.3). We report that single seed rather than spend further compute on a structural limitation of the port. CODA-Prompt reads AA 0.0 on all three CIL-CORD seeds with its default pool of 10 components; enlarging the pool to 50 (10 per session) lets some sessions be learned (diagonal 79.5/40.6/0.0/18.9/36.8) but the final row is $[0, 0, 0, 0, 36.8]$, AA 7.4, with one session never learned, so the cell stays a null. These are properties of the methods' O-tag and static-pool assumptions under five class-adding sessions, not of the sessions themselves: the same sessions are learned by naive training (AA 19.0) and by joint training (33.3), both under legacy scoring.

C. Figures and tables moved from the main paper

The CoLaR schematic (Fig. 1), the displacement profile (Fig. 2), the remedy-benefit chart (Fig. 3), the image-side comparison table (Tab. 4), the falsification ledger (Tab. 5), the memory-content ladder (Fig. 4) and two side results specified in advance (lexical routing of replay, public-proxy replay).

Routing which stored document to replay (null). The OCR vocabulary of a document is set by the input and does not drift with training, which suggests a drift-free sampler: a softmax over term-frequency signature scores $s_j = \sigma(x)^\top \sigma_j$ in three directions (*near*, *far*, *task*), uniform as the zero-point. The seed-42 reading was promising (far $+1.8$ over uniform 86.4 at d=50), so we fixed that single arm in advance for confirmation. Across three seeds the gap collapses to $+1.8/+0.4/-0.6$, mean $+0.5$ pp with a sign flip, inside the uniform-sampling spread: lexical routing is a null, and the CoLaR results do not depend on it.

Rebalancing and reading the memory (nulls). On DIL with LayoutLMv3 (CoLaR, k=4, r=128, 50 documents per task; base 87.6 on seed 42), three changes to how the memory is filled or weighted leave AA where it was: a class-balanced replay loss gives 87.1, coverage-based (k-centre) document selection 86.8 and entity-weighted compression 87.1 (seed 42 each). Each trades FUNSD retention for SROIE retention at about 1:1, so the changes redistribute retention between tasks rather than add to it. A three-seed check finds no lever above the base (class-balanced 87.8,

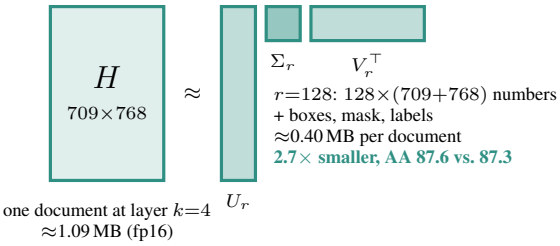


Figure 1. **CoLaR stores whole documents, compressed one at a time.** The pooled layer-k feature space is near full-rank (rank 16 keeps 43% of the variance), but a single document's activations are low-rank: its tokens share one layout and one vocabulary. Sizes for one padded document (512 text and 197 image tokens); AA on DIL, LayoutLMv3, seed 42, against uncompressed latent replay. Truncated SVD per document keeps every token's feature at its own position with its own label, which is the binding that every marginal summary destroys (Fig. 4).

Table 4. **Fixed comparison set on image CIL** (ViT-B/16, ten sessions, AA, mean ± s.d., seeds 42/7/123): the same references (naive, joint training) and baselines as the comparison table of the main paper, with SLCA and RanPAC in the slots CPFD and IS3 take on documents. ER/DER++/SLCA keep 2,000 exemplars; the uncompressed latent memory and CoLaR keep 500/class. Ridge is a diagnostic readout refit, not a method.

Method	CIFAR-100	ImageNet-R
Naive	11.7 ± 1.1	9.3 ± 0.3
LwF	15.7 ± 0.6	12.7 ± 1.4
ER (2,000 exemplars)	76.2 ± 1.2	54.8 ± 0.7
DER++ (2,000 exemplars)	76.5 ± 1.5	52.8 ± 1.1
CODA-Prompt	15.2 ± 1.8	8.4 ± 0.7
CL-LoRA (2025)	14.1 ± 1.2	11.4 ± 0.8
SLCA (2023)	87.5 ± 0.2	65.0 ± 0.3
RanPAC (2023)	85.3 ± 0.1	58.5 ± 0.3
LCA (2026)	77.2 ± 3.2	54.7 ± 3.6
Latent memory, 500/class (ours, uncompressed)	80.0 ± 0.7	55.1 ± 1.3
CoLaR r=128, 500/class (ours)	79.5 ± 1.0	55.3 ± 1.4
Ridge readout on the latent-memory trunk (probe, seed 42)	88.3	72.3
Joint training (reference)	89.5 ± 0.4	77.6 ± 0.5

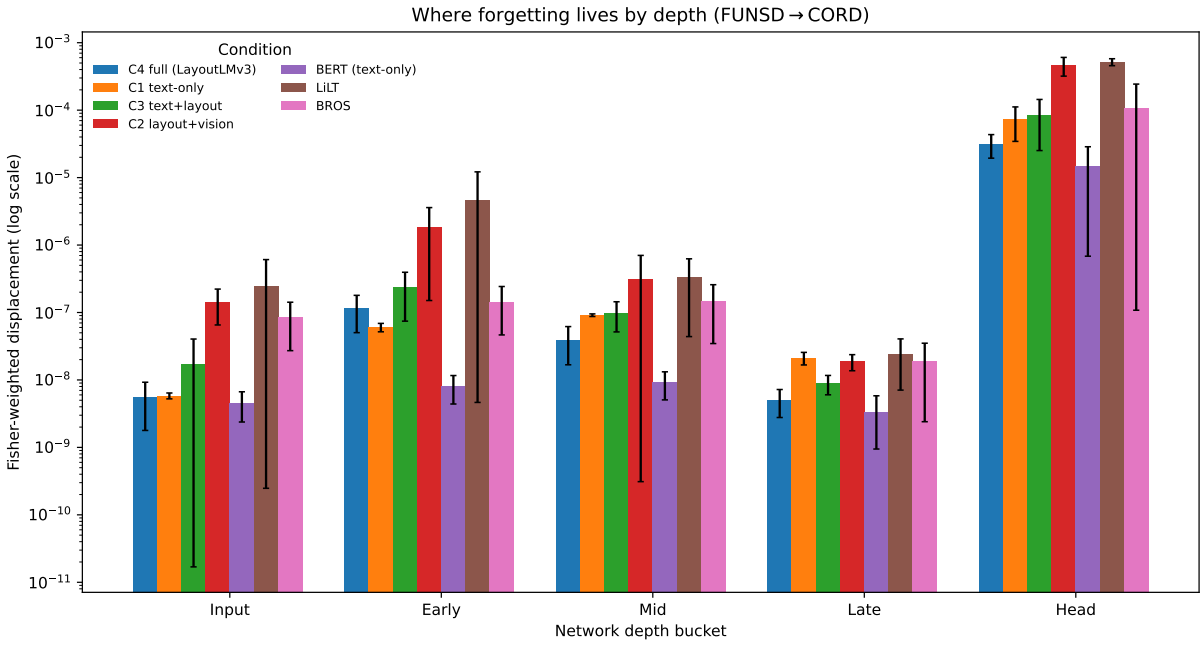


Figure 2. Where forgetting lives. Old-task-Fisher-weighted displacement D_g by depth bucket (log scale, first boundary), for the four LayoutLMv3 modality-mask conditions, text-only BERT, LiLT and BROS; mean over three seeds (two for text-only LayoutLMv3, whose third run is dead and excluded), ± 1 s.d. The classifier head's weighted displacement is two or more orders of magnitude above any encoder layer's on every backbone.

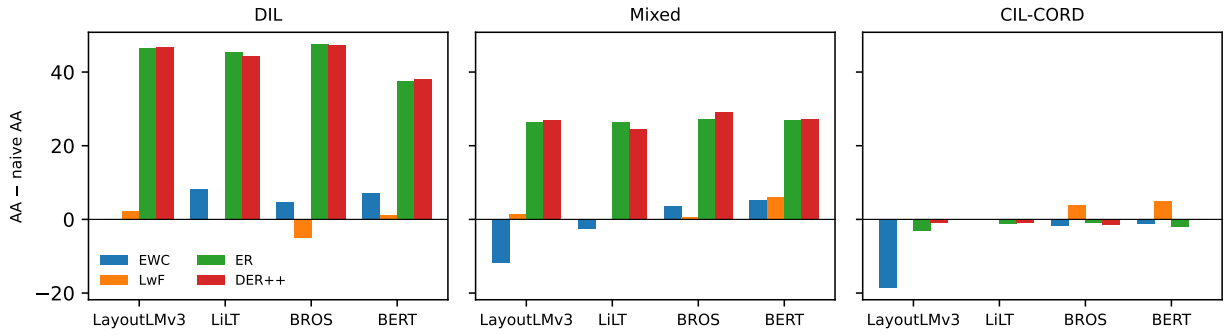


Figure 3. Remedy benefit is backbone-dependent; the phenomenon is not. Change in AA over naive for each classical remedy, per backbone and scenario (mean of three seeds). Replay's benefit is large and consistent; EWC and LwF change sign across backbones (missing bars: cells not run).

k-centre 87.0, against 87.8 ± 0.5 for the base). Two read-side additions are negative: blending a kNN readout over the stored latents into the head is inert or harmful, and metaplastic consolidation of the readout lowers AA from 85.9 to 79.5 at its own operating point.

A bounded positive on the storage side. Replaying pseudo-labelled *public* documents through the same mechanism (zero private bytes) recovers dense-label tasks (FUNSD 76.1) but not one the proxy pool does not cover (SROIE ~ 7), topping out at 58.9: coverage-limited transfer. Budget-adaptive allocation across tasks moves along a Pareto frontier, not beyond it (Appendix B).

D. Drift analysis details and legacy-scored tables

Displacement by depth, with and without the head. For a checkpoint pair $(\theta^{(t-1)}, \theta^{(t)})$ and a parameter group g we compute the old-task-Fisher-weighted displacement $D_g = \sum_{\phi \in g} F_\phi^{(t-1)} (\theta_\phi^{(t)} - \theta_\phi^{(t-1)})^2$. Tab. 6 reports it at

Table 5. Falsification ledger, DIL/LayoutLMv3 (naive 41.3, replay ≈ 88). Instance names are ours and belong to the family in the first column; LexSlot selects head slots by OCR-vocabulary overlap, and CoLaR-R is CoLaR with the stored labels refreshed. Each family is the strongest instance we could build of its mechanism; the failure reason is from a per-boundary root-cause analysis, not inferred from the score. †run at a frozen-backbone, short-schedule (3–5 epoch) operating point whose own naive level is given in the cell. ‡CIL-CORD and Mixed, LayoutLMv3, three seeds each, legacy scoring; pseudo = refresh with the model's own predictions, oracle = refresh with the true labels (not deployable).

Family	Instance	Best AA	Why it fails
Weight merging	TIES/Fisher head merge†	21.6 (no-merge 21.9)	merge of head vectors cancels; near no-op
Merge + realignment	LCA (ported)	43.9±2.5	backbone merge halves BWT; Gaussian realignment kills acquisition
Subspace transfer	HGT / CUBER (ported)†	39.8 (own naive 38.7); HGT ≤ 25.3	protected subspace grows; head unconstrained; α is a pure stability–plasticity dial
Parametric slots	LexSlot (no buffer)	42.2	isolated head slots cannot re-ground shared head
Input-anchored memory	Ledger†	FUNSD 21 / SROIE 23	own-vs-cumulative gap is training-free
Feature-Gaussian replay	LexMem v3b (positive control)	66.0	partial: recovers dense-label tasks only
	LexMem v5 (+relational)	63.2; mid-task 13.0	marginal summary breaks doc consistency
Label refresh (latent memory)	CoLaR-R pseudo (CIL-CORD / Mixed, 3 seeds each)‡	20.0 / 59.5 (plain 20.4 / 58.8; oracle 26.3 / 64.4)	teacher confident on ≈ 0 stale-O tokens; pseudo refresh is a no-op on both scenarios (oracle refresh +5.9 / +5.6 shows the slot is reachable only with true labels)

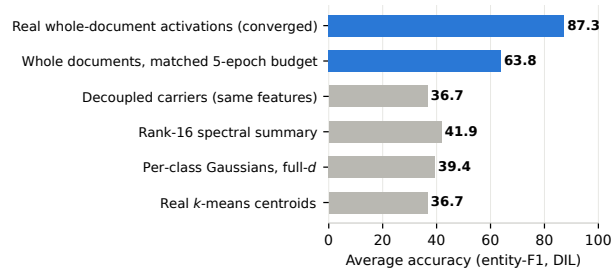


Figure 4. **What the memory must hold** (DIL, LayoutLMv3, layer k=8, AA, single-seed diagnostics). Whole documents retain; summaries of the same store fall to the naive level (41.3), and the identical features on decoupled layout carriers lose 27 points against matched whole documents.

Table 6. Fisher-weighted displacement at the first task boundary under naive fine-tuning (mean over the n live seeds; a run whose final AA is below 1.0 is dead and excluded, which removes one text-only LayoutLMv3 seed). "Head ratio": classifier D_g over the largest encoder bucket. "Encoder share": each bucket's fraction of the encoder total with the head excluded.

Condition	n	Head ratio	Encoder share input	early	mid	late
LayoutLMv3, full	3	273×	0.03	0.70	0.24	0.03
LayoutLMv3, text+layout	3	361×	0.05	0.65	0.27	0.02
LayoutLMv3, text only	2	801×	0.03	0.34	0.51	0.12
LayoutLMv3, no text	3	247×	0.06	0.80	0.13	0.01
LiLT	3	111×	0.05	0.88	0.06	0.00
BROS	3	738×	0.21	0.37	0.37	0.05
BERT	3	1608×	0.18	0.32	0.36	0.14

the first task boundary, mean over the live seeds of each condition. The classifier's D_g is two to three orders of magnitude above the largest encoder bucket on every condition (111× to 1608×). This ratio is inflated by construction: the empirical Fisher of a token classifier is largest at the head whatever the task (0.22 against 0.005 for the next group on LayoutLMv3), so the weighting itself favours the head. The informative part is the encoder-only share: with the head excluded, weighted displacement sits in the early and middle layers on every document encoder and is smallest in the late layers. Representational drift measured by CKA, in contrast, grows with depth because it accumulates. The two instruments therefore agree on the classifier and disagree on the late encoder layers, and the main text claims only what both support.

What the significance tests do and do not show. Two tests accompanied an earlier version of this analysis and we report them with their limits. (i) A Mann–Whitney test of the classifier against each other parameter group (six samples per group, Bonferroni $\alpha' = 0.0042$) rejects equality against three of twelve groups ($p = 0.0022$) and not against the rest; it was run on per-group Fisher *importance*, which is head-dominant by construction, and we do not use it as evidence of where forgetting occurs. (ii) A permutation test on the normalised D_g depth profile does not reject a shared profile between LayoutLMv3 and LiLT or BROS ($p = 0.38, 0.51$); since every normalised profile is nearly one-hot at the head, this test cannot reject and is not evidence of architecture-generality. The claim that drift is shared across backbones rests on the descriptive agreement of the CKA profiles and of Tab. 6, and on the remedy ordering of Tab. 8.

Consolidation probe. A depth-targeted consolidation method (output distillation with a per-bucket penalty) scores 42.6 AA on DIL when the penalty is spread uniformly over all depth buckets, 84.7 when concentrated on the head and late layers, 86.1 head-only and 86.2 late-only (three seeds). All four arms include a 200-document head-replay buffer, and plain ER with the same buffer scores 87.9. The probe therefore shows that a uniform penalty prevents the new tasks from being learned; it does not show that a targeted penalty recovers old tasks without replay, and we do not cite it for that.

Settings of the legacy-protocol cells. Tab. 7 lists, for every group of † cells in the main paper, the epoch cap, numeric precision, memory and seed count, read from each run's log

and saved configuration. The groups differ from each other as well as from the clean protocol, and several cells mix precisions across seeds. The DIL rows of our first grid saved no configuration; for them the table gives the grid script's defaults, which could not be confirmed per run.

Legacy-scored tables. Tabs. 8 and 9 are the tables of the earlier grid. Their CIL-CORD and Mixed columns use legacy scoring, under which joint training reads 33 on CIL-CORD; they are kept for the cross-backbone comparison (the EWC sign flip of Sec. 4.1 is read from the Mixed rows of Tab. 8) and for completeness. "Independent" is the mean of one model trained per task from the pre-trained checkpoint.

Table 7. **Settings of the legacy-protocol (†) cells.** Every row early-stops on the session's test split, except joint training on CIL-CORD and Mixed, which holds out 10% of its pooled training set. Precision "fp32/fp16": seed 42 on CIL-CORD ran in fp32 (mixed precision was added to these training loops only afterwards) and every other seed in fp16 mixed precision. Seeds are given per scenario in the order of the second column. The clean protocol replaces all of this with fp32, a 10-epoch cap (100 for joint training) and early stopping on 10% held out from each session's training data. ‡LwF: the one fp16 seed per scenario that collapsed to zero is excluded (CIL-CORD seed 123, Mixed seed 7), leaving 2 and 1 seeds. *The first DIL grid (naive, EWC, LwF, ER, DER++, CODA-Prompt, CL-LoRA, joint training; all four backbones) saved no configuration; the values are the grid script's defaults and could not be confirmed per run (that script's default batch size was 16, against 2 for every later run), and joint training early-stopped on the pooled test sets.

Row	Scenarios	Epoch cap	Precision	Memory	Seeds
Naive, joint training	CIL-CORD, Mixed	100	fp16	none	3, 2
ER, DER++	CIL-CORD, Mixed	100	fp32/fp16	200 samples (reservoir)	3, 2
LwF	CIL-CORD, Mixed	100	fp32/fp16	none	3, 2‡
CPFD	CIL-CORD, Mixed	100	fp32/fp16	none	2, 2
IS3	CIL-CORD, Mixed	100	fp32	none	2, 2
CODA-Prompt, CL-LoRA	CIL-CORD, Mixed	30	fp32/fp16	none	2, 2
LCA	CIL-CORD, Mixed	10	fp32	class Gaussians	2, 2
CoLaR	CIL-CORD, Mixed	10	fp16	50 docs/session	2, 2
CoLaR-SBR	CIL-CORD	10	fp32	50 docs/session	3
CoLaR-SBR	Mixed	10	fp16	50 docs/session	2
SBR without memory	CIL-CORD	10	fp32	none	1
CoLaR-R, sliced replay loss	CIL-CORD	10	fp16	50 docs/session	2
CoLaR, true-label memory refresh	CIL-CORD	10	fp32	50 docs/session	1
SBR-Recall	CIL-CORD, Mixed	10	fp32	none, 50 docs/session	3, 1
CoLaR-SBR + LexSlot	CIL-CORD, Mixed, DIL	10	fp32	50 docs/session	3, 3, 3
CoLaR, CoLaR-SBR, CPFD, IS3, LCA	DIL	10	fp32	CoLaR rows: 50 docs/task	3 (IS3: 1)
First-grid rows*	DIL	100 (CODA-Prompt 30)	unknown	ER, DER++: 200 samples	3

Table 8. Generality grid (legacy scoring on CIL-CORD and Mixed, test-split early stopping): average accuracy (mean ± s.d., three seeds) of the six classical strategies on four backbones and three scenarios. *Indep.* is a per-dataset reference (one model trained independently on each dataset, three-seed macro over FUNSD/SROIE/CORD). *Joint* is joint training on the pooled data (one model, all three datasets), also a reference. Independent is only meaningful on DIL, where the task stream is three distinct datasets; CIL-CORD is one dataset split by label and Mixed is a composed stream, so Indep. is undefined there and marked N/A. Grey cells did not fit the 6 GB GPU. On DIL replay comes close to joint training on every backbone; regularisation and distillation stay within ten points of naive training on average; on CIL-CORD every method stays below joint training under this scoring, and only LwF (on BROS and BERT) rises above naive training.

Scenario	Backbone	Naive	Indep.	Joint	EWC	LwF	ER	DER++
DIL	LayoutLMv3	41.3±1.5	88.8±0.6	88.7±1.1	41.2±1.1	43.7±1.1	87.9±1.4	88.0±0.9
	LiLT	40.1±0.5	85.3±0.8	83.8±1.2	48.4±5.9	–	85.7±0.5	84.4±0.3
	BROS	40.3±1.8	88.8±0.1	87.5±1.3	45.0±4.9	35.1±3.1	87.8±0.4	87.5±1.1
	BERT	39.4±0.3	77.3±0.5	77.2±0.4	46.6±3.0	40.5±0.1	77.1±0.9	77.4±0.6
CIL-CORD	LayoutLMv3	19.0±0.2	N/A	33.3±0.3	0.5±0.6	18.8±0.2	16.0±1.4	18.2±0.3
	LiLT	19.1±0.3	N/A	32.4±0.4	–	–	17.9±0.5	18.0±1.0
	BROS	19.2±0.1	N/A	32.9±0.0	17.3±0.8	23.1±1.2	18.3±1.5	17.7±1.1
	BERT	19.3±0.1	N/A	32.4±0.2	17.9±0.3	24.3±0.3	17.1±1.6	19.2±1.9
Mixed	LayoutLMv3	33.9±0.3	N/A	67.5±0.8	22.0±9.4	35.2±1.2	60.3±5.3	60.8±2.7
	LiLT	30.7±0.7	N/A	61.8±0.6	28.1±3.0	–	57.0±1.4	55.1±2.9
	BROS	33.9±0.4	N/A	65.3±2.1	37.4±3.3	34.4±0.5	61.1±1.4	63.0±0.4
	BERT	23.4±0.3	N/A	55.0±0.6	28.6±0.9	29.4±4.2	50.2±0.6	50.7±1.3

Table 9. **Fixed comparison set on LayoutLMv3 under legacy scoring** (AA, mean ± s.d. over seeds 42/7/123, legacy protocol): naive and joint training as references and the published baselines (CPFD and IS3 in the slots SLCA and RanPAC take on images), then our rows; the image-side table is Tab. 4. *CoLaR-R, sliced*: CoLaR with the replay loss restricted to the classes alive when each document was stored. Its CIL-CORD and Mixed columns are runs of the earlier grid, not the † cells of the main paper: CoLaR, CoLaR-R, LexSlot, CPFD, IS3 and LCA ran in fp32 with a 10-epoch cap; the other rows saved no configuration (grid default cap 100, 30 for CODA-Prompt). The DIL column is the main paper's. LexSlot is the 200-exemplar slot-memory recipe. *did not converge on these sessions, which naive (19.0) and joint training (33.3) do learn (Appendix B, "Null cells among the fixed baselines"). "–": not applicable; IS3 on DIL is a single seed (fixed-label port limitation, Appendix B).

Method	**CIL-CORD**	**DIL**	**Mixed**
Naive	18.98 ± 0.17	41.31 ± 1.54	33.85 ± 0.33
LwF	18.82 ± 0.15	43.66 ± 1.05	35.23 ± 1.17
ER	15.98 ± 1.36	87.85 ± 1.42	60.32 ± 5.25
DER++	18.15 ± 0.30	**88.02 ± 0.88**	**60.82 ± 2.74**
CODA-Prompt	null*	32.51 ± 1.23	8.11 ± 1.52
CL-LoRA (2025)	15.17 ± 3.14	41.14 ± 0.28	31.37 ± 0.63
CPFD (2023)	null*	6.15 ± 7.29	null*
IS3 (2024)	null*	38.29 (1 seed)	30.10 ± 1.18
LCA (2026)	8.15 ± 2.59	43.87 ± 2.45	19.01 ± 3.68
Text-only comparator (no layout or image):			
BERT-base	19.26 ± 0.11	39.42 ± 0.26	23.37 ± 0.27
CoLaR (ours)	20.38 ± 1.34	87.75 ± 0.48	58.75 ± 1.10
CoLaR-R, sliced (ours)	25.38 ± 2.38	–	61.30 ± 1.91
LexSlot (ours)	19.28 ± 0.12	87.33 ± 1.03	35.49 ± 1.39
Joint training (reference)	33.32 ± 0.33	88.73 ± 1.13	67.51 ± 0.82
Independent (one model per task)	–	88.79 ± 0.55	–