


Per-layer representational drift (linear CKA between consecutive checkpoints, C4)



Finding: a monotonic depth gradient — input encoders do not drift (CKA = 1.0), while drift rises with depth and peaks at the last encoder layer and the classifier. Forgetting is a late-layer + head phenomenon.