


## Naive sequential fine-tuning: per-task F1 across the sequence (pilot, C4, seed 42)



*Finding: each task is learned to a strong F1 (bold markers: 88 / 93 / 80.5) and then collapses to ~2–4 F1 once the next task is trained. These curves are the rows of the matrix in Fig.~6.1 read left-to-right.*