


Performance matrix $R[i,j]$ under naive sequential fine-tuning (C4, seed 42)



Finding: forgetting is near-total within one task step

Each task is learned well in isolation (diagonal 80–93 F1), but after one further task the earlier task collapses: FUNSD falls 88.0 $\rightarrow$ 1.8 (–86 F1) once CORD is learned. The text+layout condition C3 shows the identical pattern (88.7 $\rightarrow$ 0.3 $\rightarrow$ 1.7).



retained / just-learned



forgotten



not yet trained