

Performance matrix $R[i, j]$ — naive sequential fine-tuning
diagonal stays high, retention collapses to $\sim 2\text{--}4$ F1

