


Where forgetting lives in LayoutLMv3 — synthesis of the per-component diagnosis



Preliminary — pilot diagnosis (C4, naive sequential, 3 seeds).

Finding

Forgetting concentrates in the classifier head and late encoder layers; the input encoders are stable. It is **not** a fusion-specific effect — the full multimodal model is statistically indistinguishable from its modality ablations (GATE-A). Text is the load-bearing stream.

Hypothesis (open)

Because the forgetting locus is output-facing, remedies that protect or re-expose the head — e.g. experience replay — should be most effective. §6.2 confirms replay essentially eliminates forgetting on DIL. A mechanism that targets the head / late layers specifically (DocCL) is the open extension.